Elastic Load Balancing & Auto Scaling

1) Create Launch Template

- Go to **EC2 → Launch Templates**
- Click **"Create launch template"**
- Configure the following:

Template name: web-server-template

AMI: Amazon Linux 2 (ami-0abcdef1234567890)

Instance type: t2.micro

Key pair: Select existing or create new

Security groups: Create new with HTTP (80) and SSH (22)

User Data Script

Add this script in Advanced details → User data:

```
#!/bin/bash
yum update -y
yum install -y httpd
systemctl start httpd
systemctl enable httpd
echo "<h1>Web Server $(hostname -f)</h1>" >
/var/www/html/index.html
```

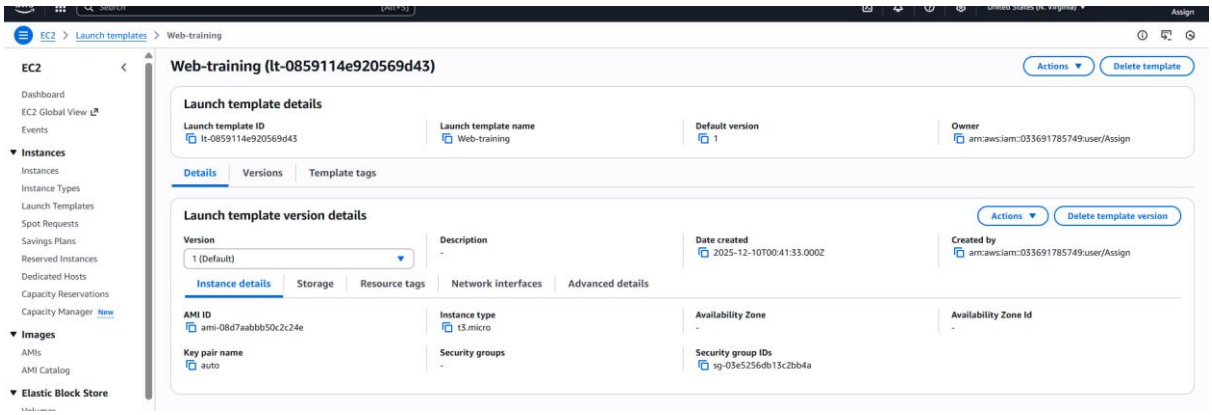

2) Create Auto Scaling Group

1. **Go to EC2 → Auto Scaling Groups**

2. **Click "Create Auto Scaling group"**
3. **Select your launch template**

   Auto Scaling group name: webtrain

   Launch template: web-server-template

   VPC: Default VPC

   Subnets: Select 2+ subnets in different AZs

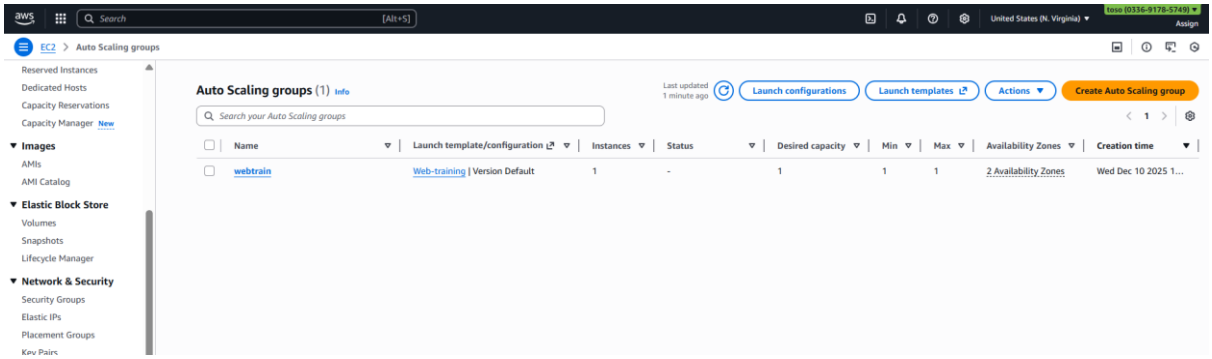

3) Create Application Load Balancer

- Go to **EC2 → Load Balancers**
- Click **"Create Load Balancer"**
- Choose **"Application Load Balancer"**

Name: webtrain
Scheme: Internet-facing
IP address type: IPv4
VPC: Default VPC
Mappings: Select 2+ AZs with public subnets

Target Group Configuration

- **Target type: Instances**
- **Protocol: HTTP**
- **Port: 80**
- **Health check path: /**
- **Health check interval: 30 seconds**

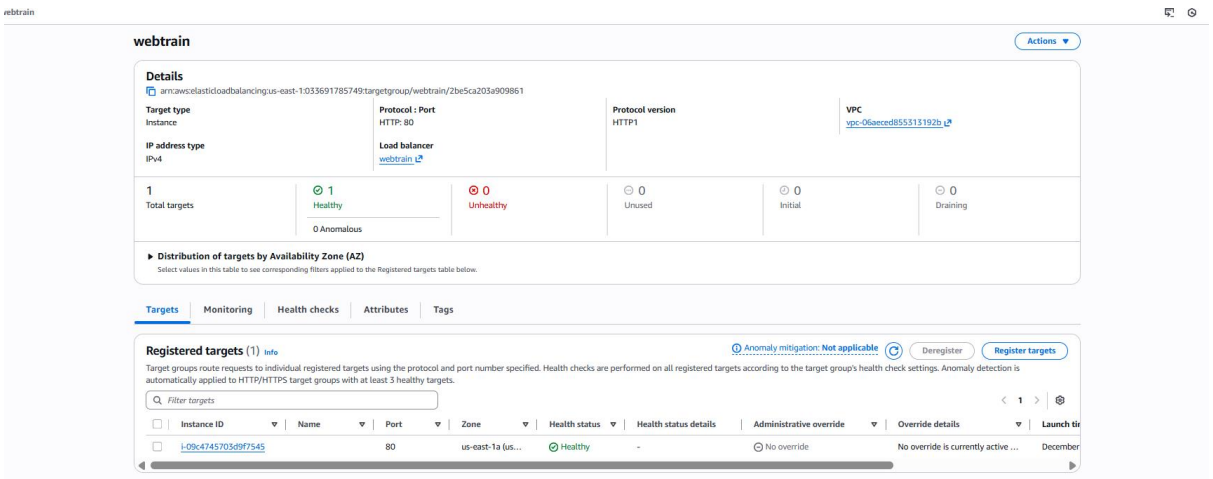

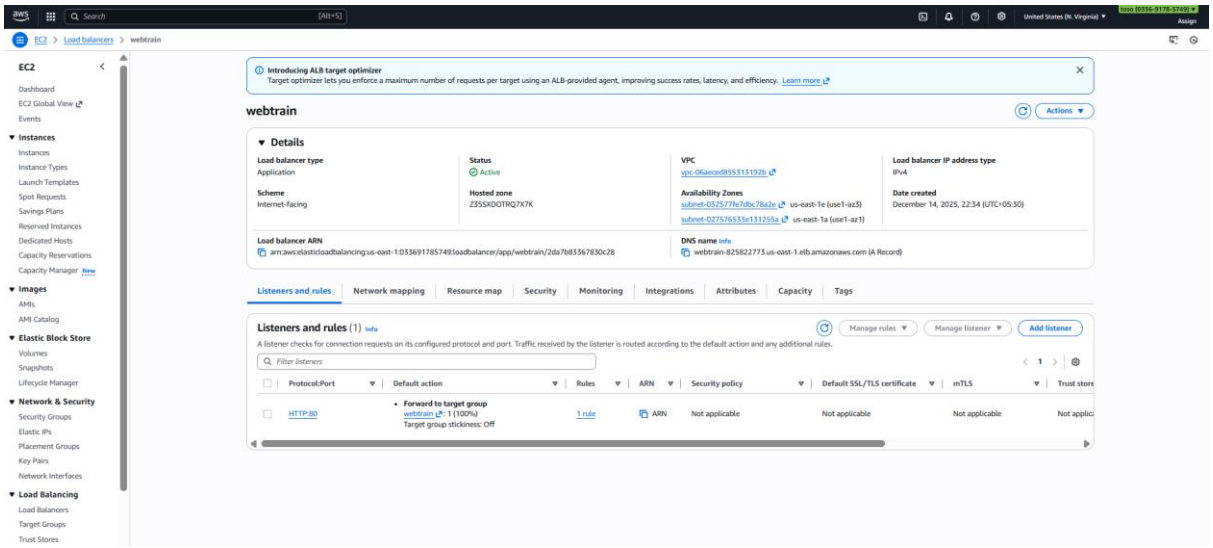

4) Attach Load Balancer to Auto Scaling Group

- Go back to **Auto Scaling Groups**
- Select your ASG: **web-server-asg**
- Go to **Details** tab
- Click **"Edit"** in Load balancing section

Load balancing: Enable

Target groups: Select your target group

Health check type: ELB

Health check grace period: 300 seconds

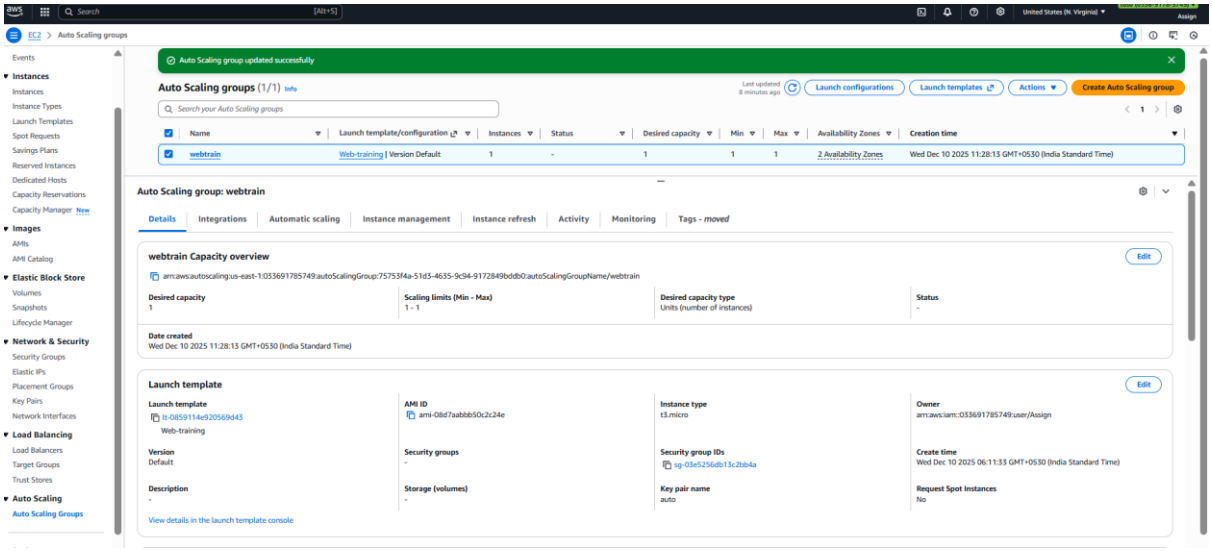

5) Configure Dynamic Scaling Policies

6) n ASG, go to **Automatic scaling** tab

7) Click **"Create dynamic scaling policy"**

Policy type: Target tracking scaling

Metric type: Average CPU Utilization

Target value: 70%

Instance warmup: 300 seconds

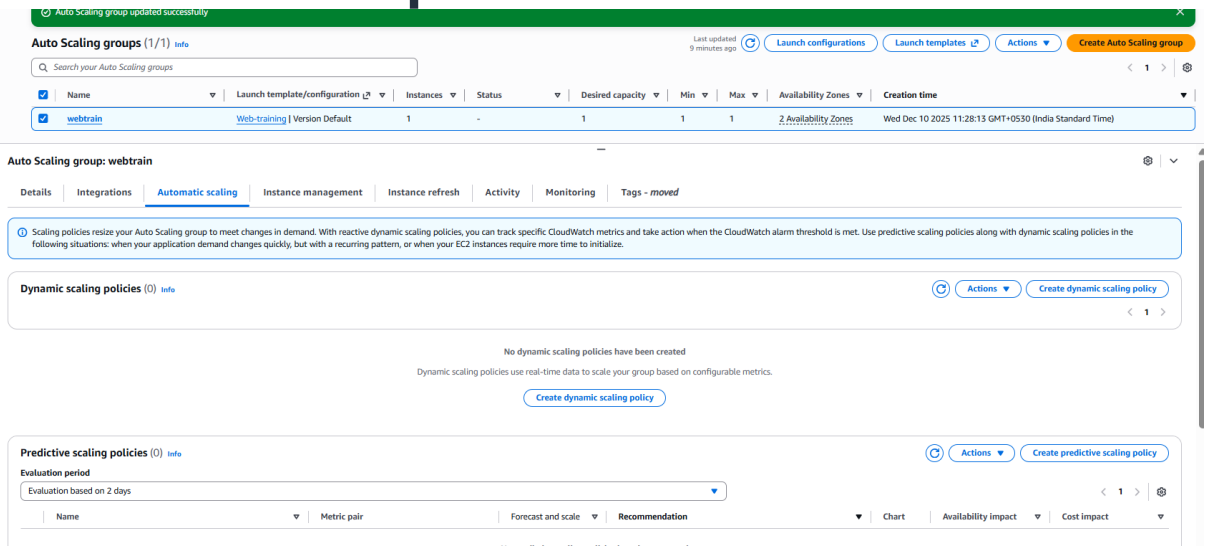

Create dynamic scaling policy

Policy type

Target tracking scaling

Scaling policy name

Target Tracking Policy

Metric type | Info

Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Average CPU utilization

Target value

70

Instance warmup | Info

300 seconds

☐ Disable scale in to create only a scale-out policy